

# **COMP2211: Networks and Systems**

## **Networks: Application Layer**

Donald Sturgeon

[donald.j.sturgeon@durham.ac.uk](mailto:donald.j.sturgeon@durham.ac.uk)

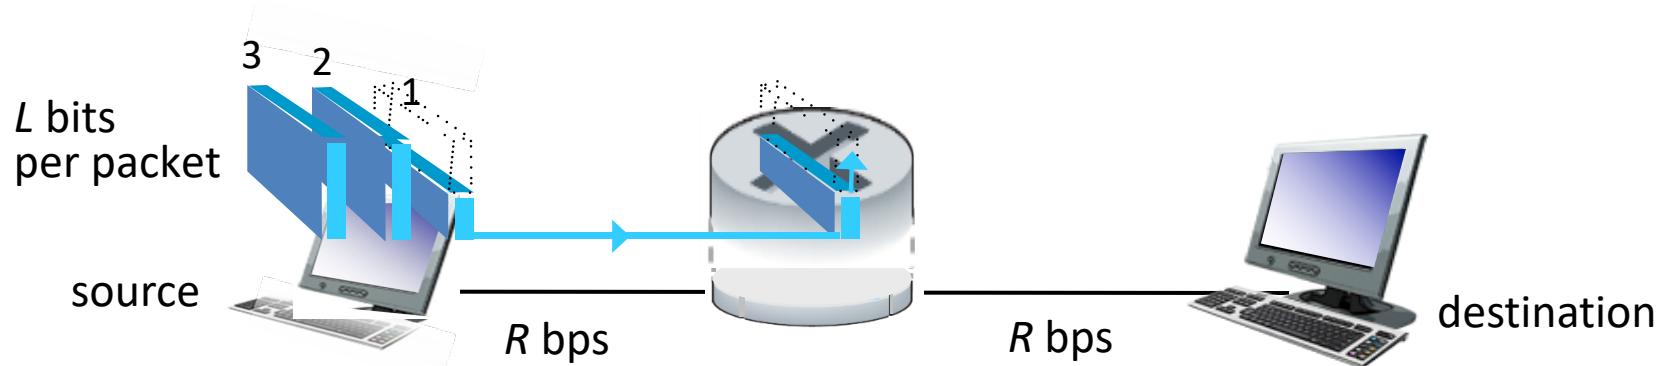
Department of Computer Science

# Previously

---

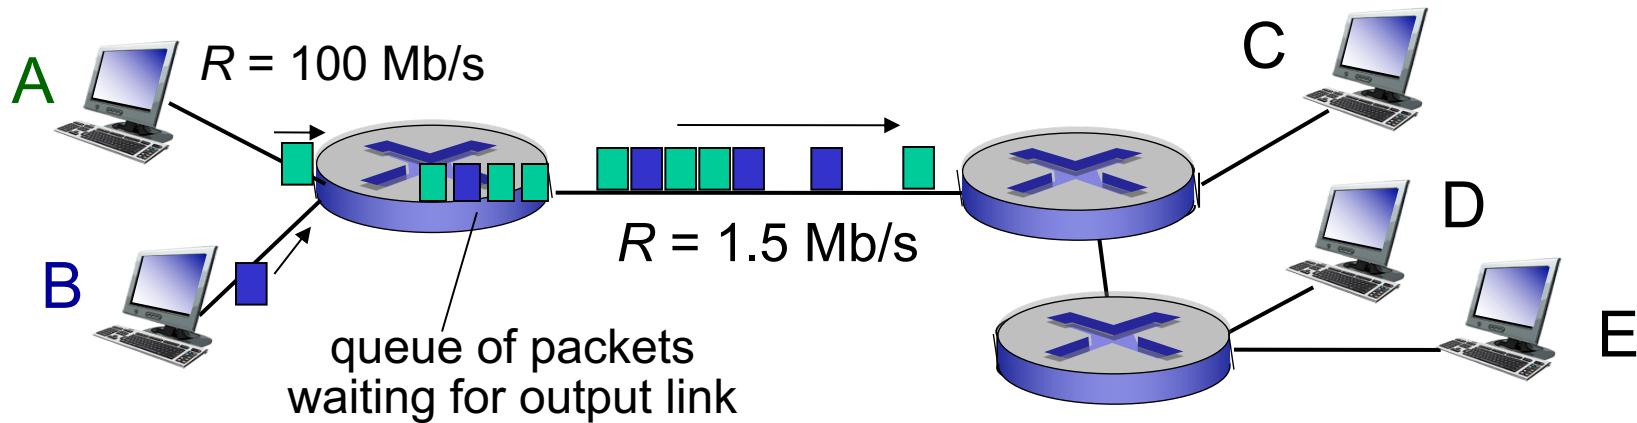
- Network core
  - Packet switching
  - Circuit switching
- Delay and loss in networks
- Protocol layers

# Packet-switching: store-and-forward



- ***Store and forward:*** entire packet must arrive at router before it can be transmitted on next link
- Takes  $L/R$  seconds to transmit (push out)  $L$ -bit packet into link at  $R$  bps

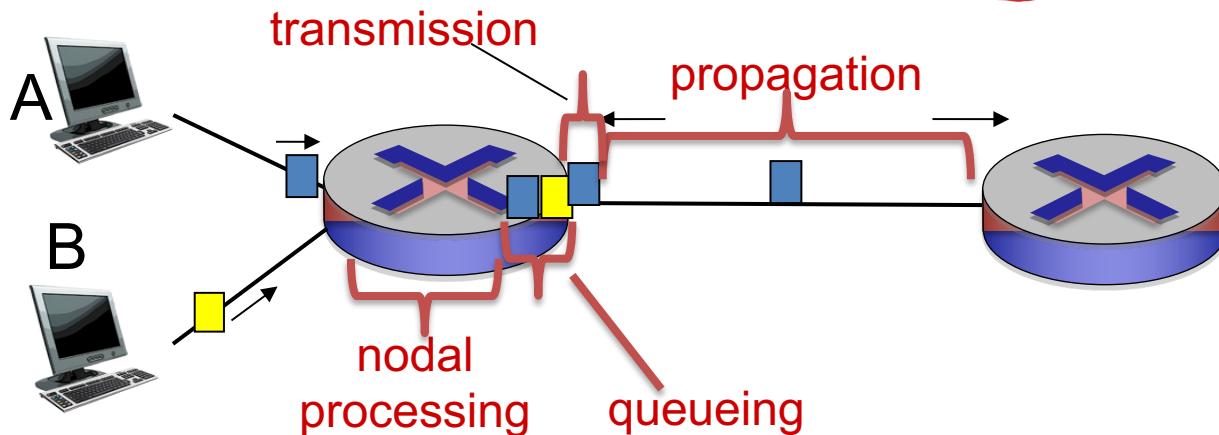
# Packet Switching: queuing delay, loss



## Queuing and loss:

- If arrival rate (in bits) to link exceeds transmission rate of link for a period of time:
  - Packets will queue, wait to be transmitted on link.
  - Packets can be dropped (lost) if memory (buffer) fills up.

# Four sources of packet delay



$d_{\text{proc}}$ : nodal processing

- Check bit errors
- Determine output link
- Typically < msec

$d_{\text{trans}}$ : transmission delay:

- $L$ : packet length (bits)
- $R$ : link bandwidth (bps)
- $d_{\text{trans}} = L/R$

$d_{\text{queue}}$ : queueing delay

- Time waiting at output link for transmission
- Depends on congestion level of router

$d_{\text{prop}}$ : propagation delay:

- $d$ : length of physical link
- $s$ : propagation speed
- $d_{\text{prop}} = d/s$

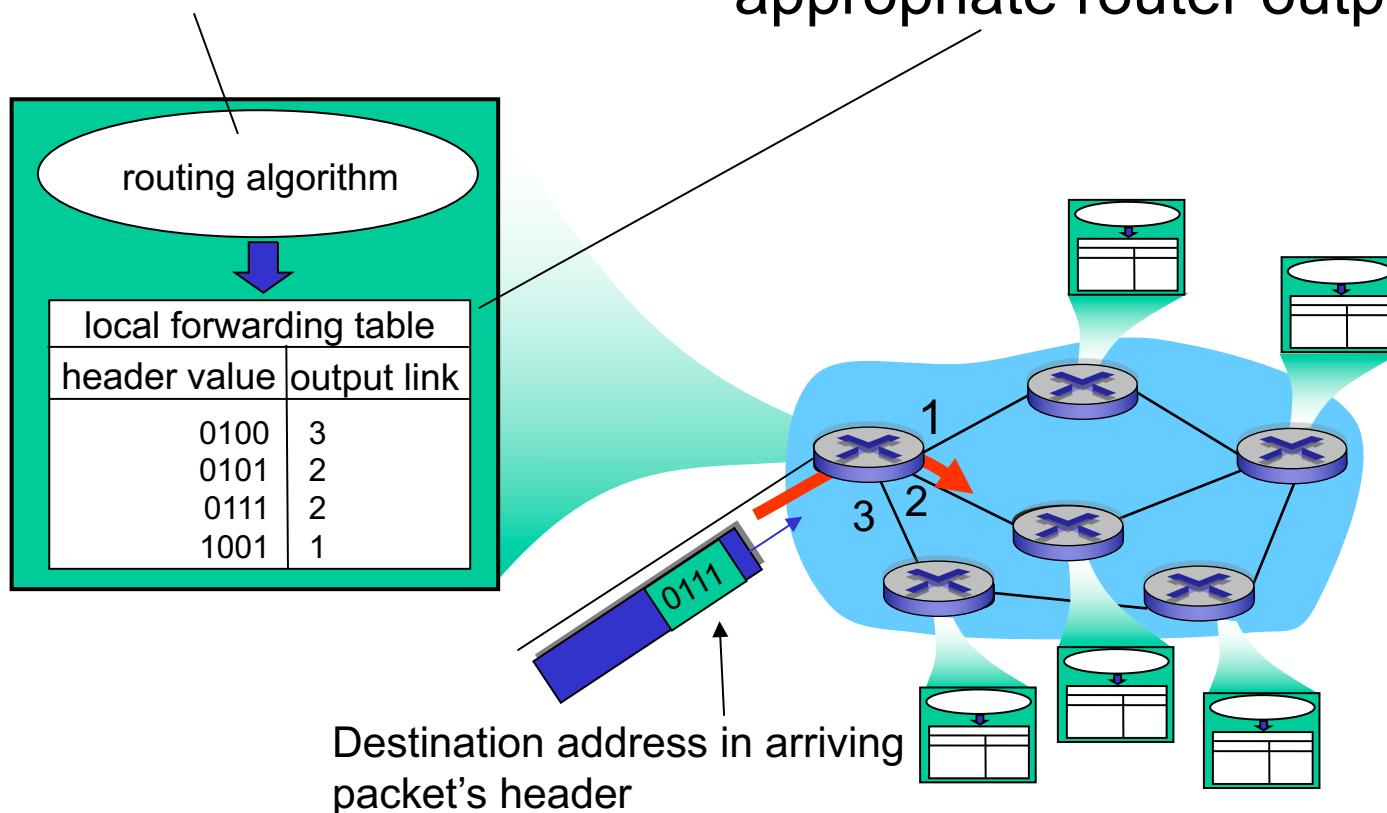
$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

# Two key network-core functions

**Routing:** determines source-destination route taken by packets

- *routing algorithms*

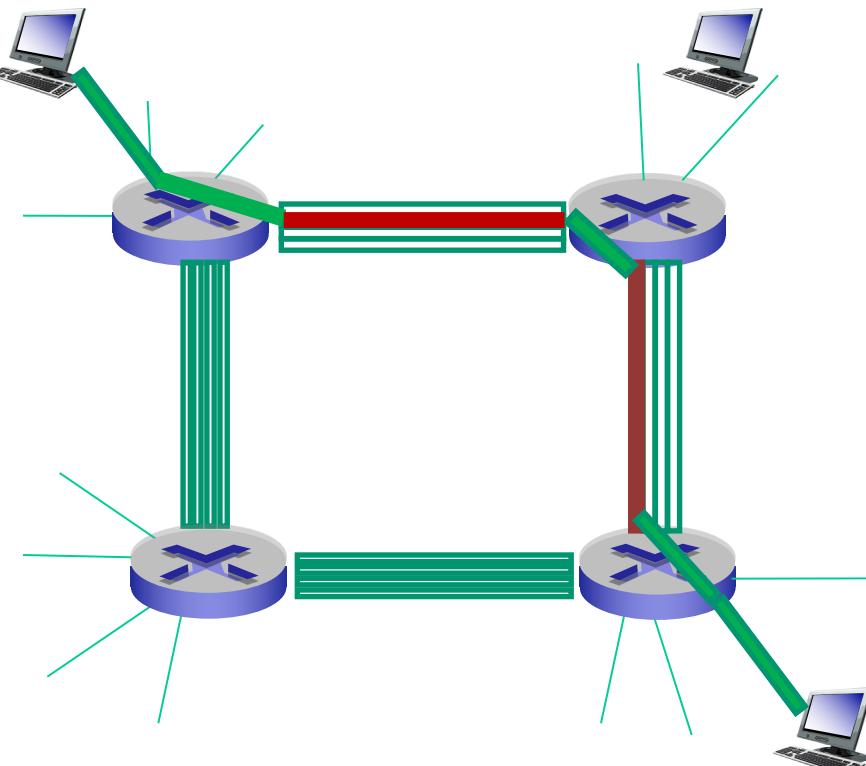
**Forwarding:** move packets from router's input to appropriate router output



# Alternative core: circuit switching

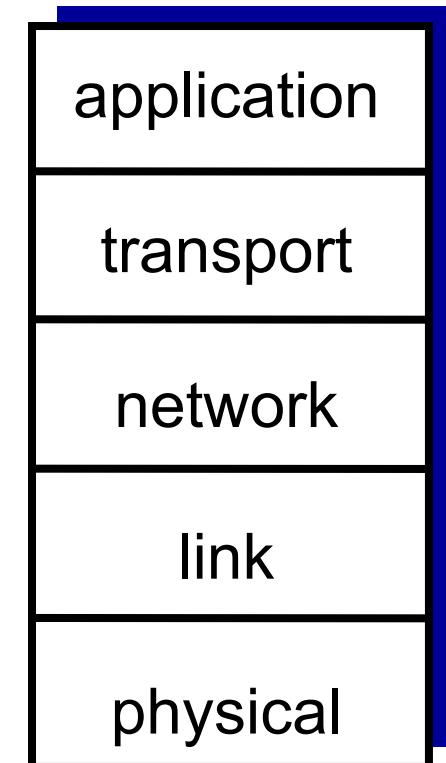
End-end resources allocated to,  
reserved for “call” between source  
& dest:

- In diagram, each link has four circuits.
  - Call gets **2<sup>nd</sup> circuit** in top link  
and **1<sup>st</sup> circuit** in right link.
- Dedicated resources: no sharing
  - Circuit-like (guaranteed) performance
- Circuit segment idle if not used by call (**no sharing**)
- Commonly used in traditional telephone networks



# Internet protocol stack

- ***Application:*** supporting network applications
  - FTP, SMTP, HTTP
- ***Transport:*** process-process data transfer
  - TCP, UDP
- ***Network:*** routing of datagrams from source to destination
  - IP, routing protocols
- ***Link:*** data transfer between neighboring network elements
  - Ethernet, 802.11 (WiFi), PPP
- ***Physical:*** bits “on the wire”



# Application Layer

---

- ◆ Network architecture:
  - ◆ Client-server architecture
  - ◆ P2P architecture
- ◆ Protocols:
  - ◆ TCP (Transmission Control Protocol)
  - ◆ UDP (User Datagram Protocol)
- ◆ Socket programming:
  - ◆ TCP
  - ◆ UDP

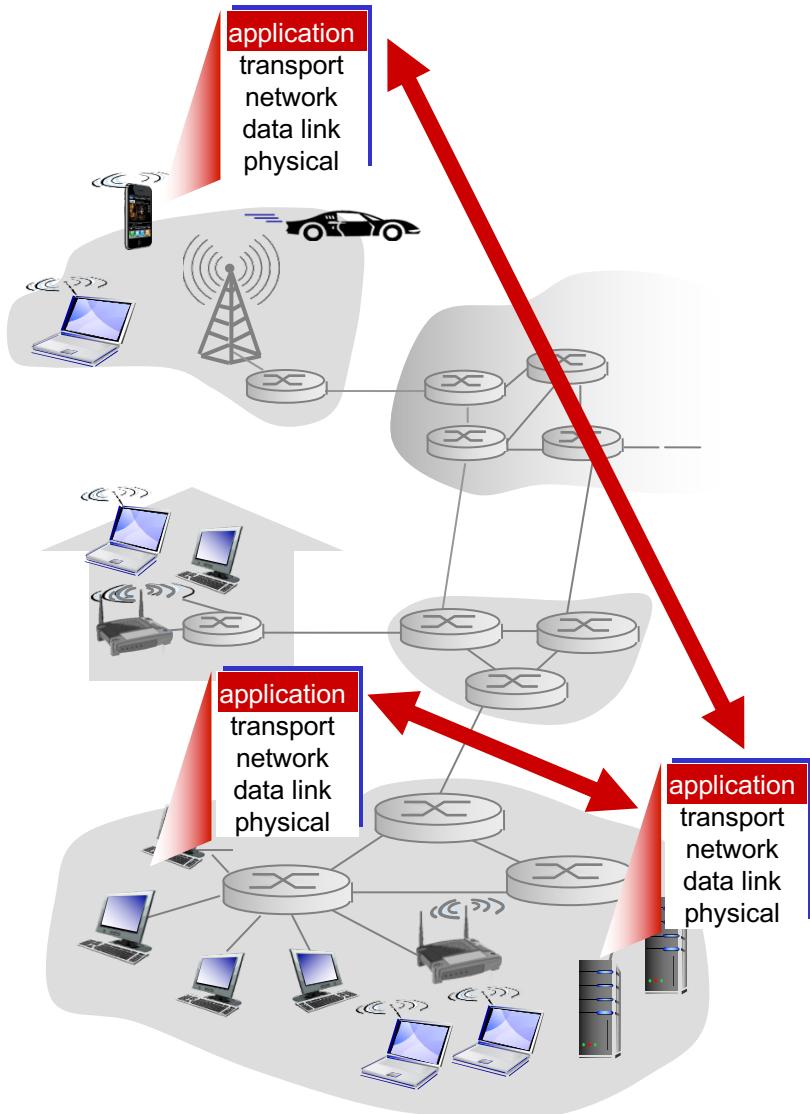
# Creating a network app

**Write programs that:**

- Run on (different) end systems
- Communicate over network
- E.g., web server software communicates with browser software

**No need to write software for network-core devices**

- Network-core devices do not run user applications
- Applications on end systems allow for rapid app development, propagation

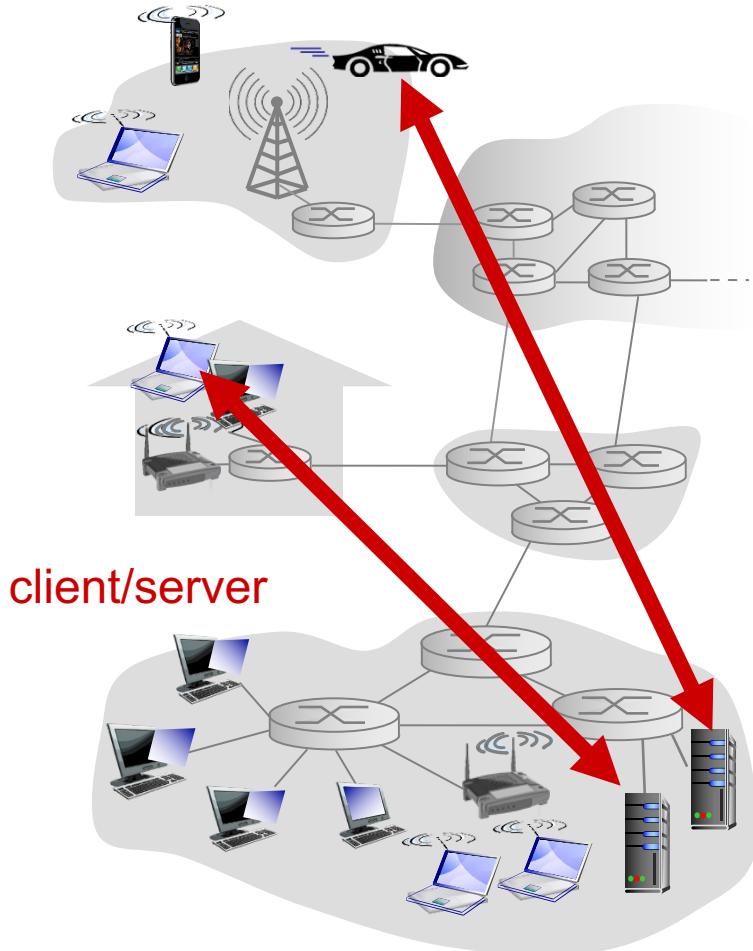


# Application architectures

Possible structures of applications:

- Client-server
- Peer-to-peer (P2P)

# Client-server architecture



## Server:

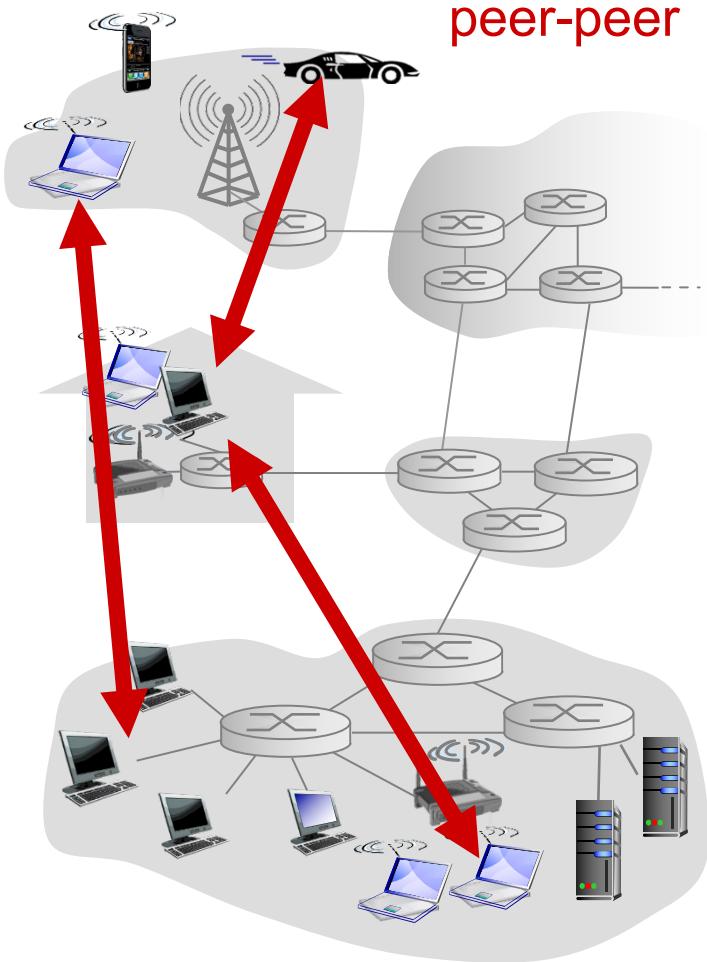
- Always-on host
- Fixed (static) IP address
- Data centres for scaling

## Clients:

- Communicate with server
- May be intermittently connected
- May have dynamic IP addresses
- Do not communicate directly with each other

# P2P architecture

- No always-on server
- Arbitrary end systems directly communicate
- Peers request service from other peers, provide service in return to other peers
  - *Self scalability* – new peers bring new service capacity, as well as new service demands
- Peers are intermittently connected and change IP addresses



peer-peer

# Road map

---

- ◆ Network architecture:
  - ◆ Client-server architecture
  - ◆ P2P architecture
- ◆ Processes:
  - ◆ TCP
  - ◆ UDP
- ◆ Web and HTTP
- ◆ Socket programming:
  - ◆ TCP
  - ◆ UDP

# Processes communicating

**Process:** program running within a host

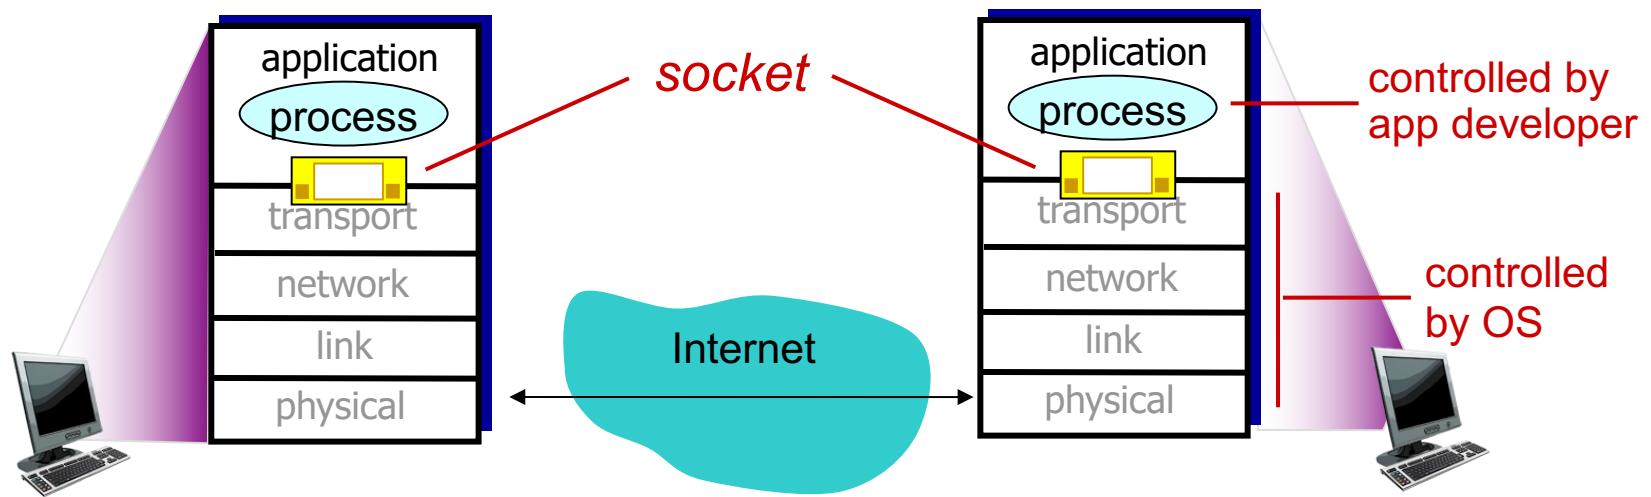
- Within same host, two processes communicate using **inter-process communication** (defined by OS)
- Processes in different hosts communicate by exchanging messages

**Socket:** a software mechanism that allows a process to create and send messages into, and receive messages from the network.

- A process is analogous to a house, and its socket is analogous to its door
- Interface between application layer and transport layer.

# Sockets

- Process sends/receives messages to/from its **socket**
- Socket analogous to door
- **Sending process shoves message out of the door**
- Sending process relies on transport infrastructure on other side of the door to deliver message to a socket at the receiving process



# What transport service does an app need?

---

## Data integrity

- Some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- Other apps (e.g., audio) can tolerate some loss

## Security

- Encryption, data integrity, ...

## Timing

Some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

# Internet transport protocols services

---

## TCP service:

- ***Connection-oriented***: setup required between client and server processes
- ***Reliable transport*** between sending and receiving process
- ***Flow control***: sender won't overwhelm receiver
- ***Full-duplex connection***: connection can send messages to each other at the same time

## UDP service:

- ***Unreliable data transfer*** between sending and receiving processes
- ***Does not provide***: reliability, flow control, congestion control, timing, security, or connection setup

# App-layer protocol defines

- Types of messages exchanged
  - E.g. request, response
- Message syntax
  - What fields in messages & how fields are delineated
- Message semantics
  - Meaning of information in fields
- Rules for when and how processes send & respond to messages

**Open protocols:**

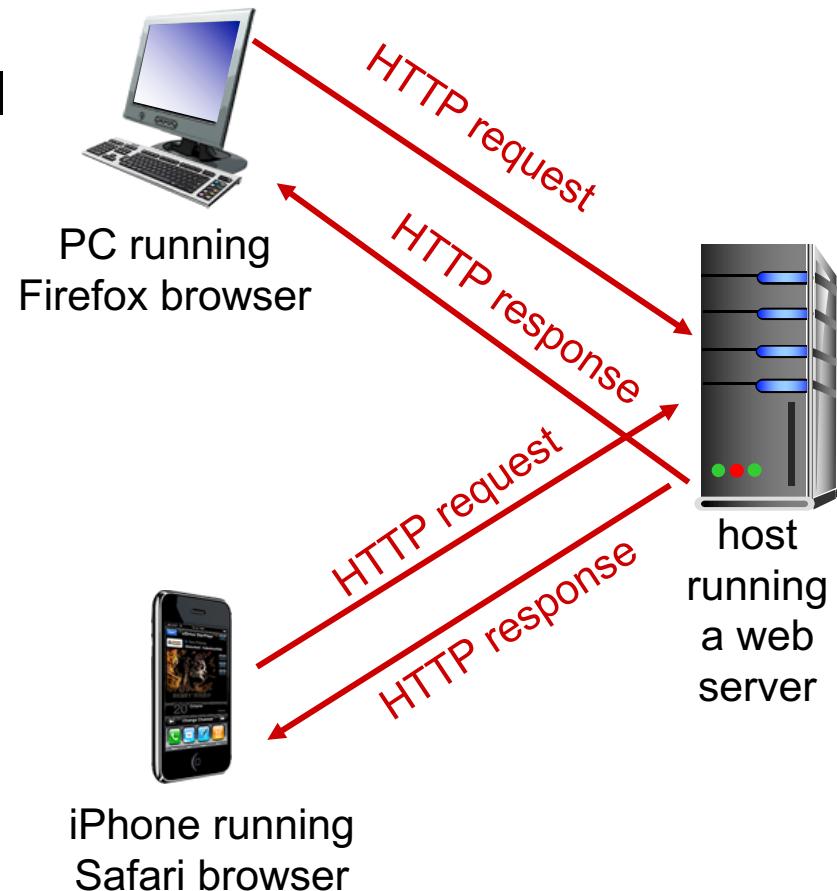
Defined in Request For  
Comments (RFC)  
Allow for interoperability,  
e.g. HTTP, SMTP

Proprietary protocols  
e.g. Skype

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- Client/server model
  - **Client:** browser that requests, receives, (using HTTP protocol) and “displays” Web objects
  - **Server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview (continued)

## Uses TCP:

- Client initiates TCP connection (creates socket) to server, port 80
- Server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## *HTTP is “stateless”*

Server maintains no information about past client requests

*aside*

### Protocols that maintain “state” are complex!

- Past history (state) must be maintained
- If server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

## *Non-persistent HTTP*

- At most one object sent over TCP connection
  - Connection then closed
- Downloading multiple objects required multiple connections

## *Persistent HTTP*

- Multiple objects can be sent over single TCP connection between client, server

# Non-persistent HTTP

Suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

Ib. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. **“accepts” connection**, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket.  
Message indicates that client wants object CS-department/home.index

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time

# Non-persistent HTTP (cont.)

time  
↓

5. HTTP client receives response message containing html file, displays html. Parsing html file, finds  $n$  referenced objects (e.g., JPEG)

4. HTTP server closes TCP connection. But, not really! (until the client received the response intact)

6. First 4 steps are repeated for each of 10 jpeg objects

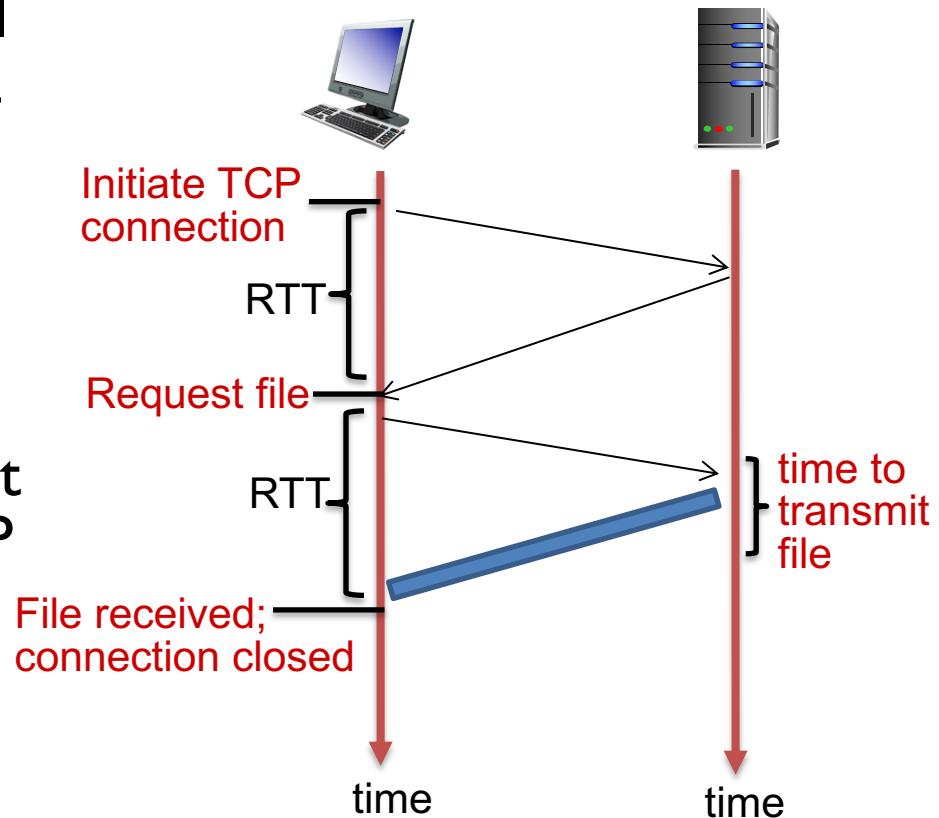
# Non-persistent HTTP: response time

**RTT (round-trip time)**

(definition): time for a small packet to travel from client to server and back

**HTTP response time:**

- One RTT to initiate TCP connection
  - One RTT for HTTP request and first few bytes of HTTP response to return
  - File transmission time
  - Non-persistent HTTP response time =  $2\text{RTT} + \text{file transmission time}$
- Incurred for each file



# Persistent HTTP

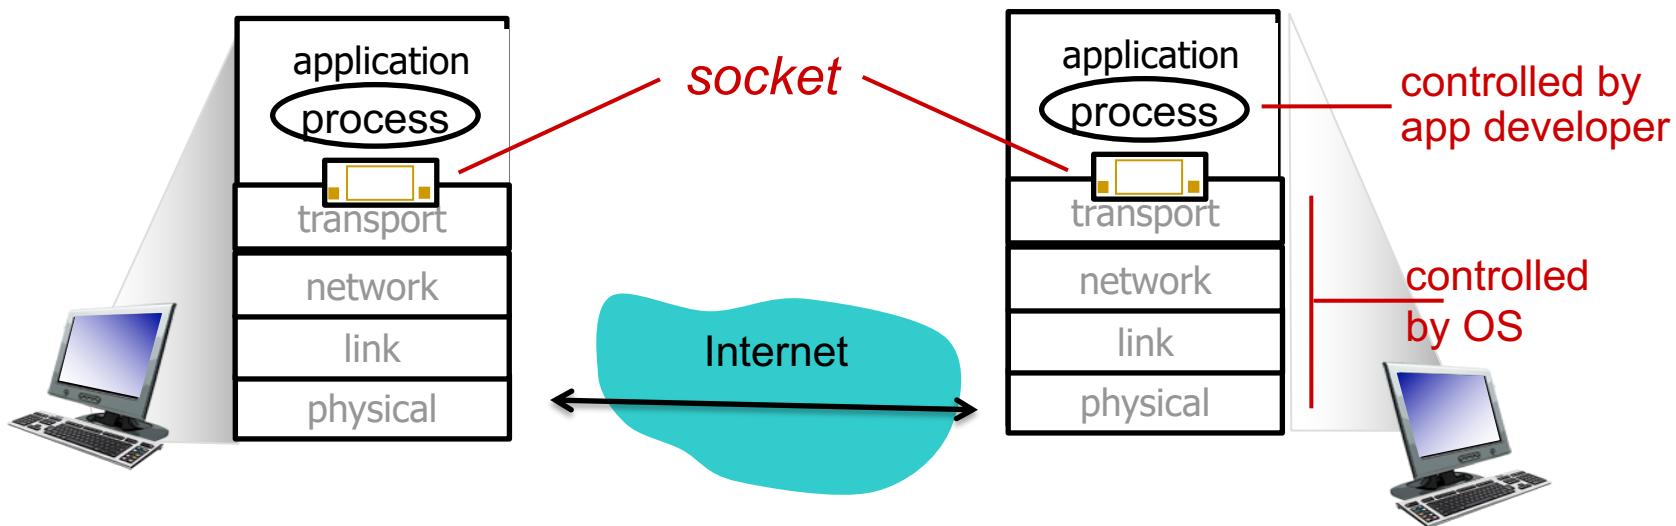
## *Persistent HTTP:*

- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server sent over open connection
- Client sends requests as soon as it encounters a referenced object
- Takes as little as one RTT + file transmission time *total*
  - Assuming connections to server already established
  - Assuming all files requested in parallel

# Socket programming

**Goal:** learn how to build client/server applications that communicate using sockets

**Socket:** door between application process and end-end-transport protocol



# Socket programming

*Two socket types for two transport services:*

- **UDP**: unreliable datagram
- **TCP**: reliable, byte stream-oriented

*Application Example:*

1. Client reads a line of characters (data) from its keyboard and sends data to server
2. Server receives the data and converts characters to uppercase
3. Server sends modified data to client
4. Client receives modified data and displays line on the screen

# Socket programming *with UDP*

---

- UDP: no “connection” between client & server
- No handshaking before sending data
- Sender explicitly attaches IP destination address and port # to each packet
- Receiver extracts sender IP address and port# from received packet
- UDP: transmitted data may be lost or received out-of-order
- Application viewpoint:
- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

# Client/server socket interaction: UDP

## Server (running on some IP)

Create socket, port = x:

```
serverSocket =  
socket(AF_INET,SOCK_DGRAM)
```

Read datagram from  
**serverSocket**

Write reply to  
**serverSocket**  
specifying  
client address,  
port number

## Client

Create socket:

```
clientSocket =  
socket(AF_INET,SOCK_DGRAM)
```

Create datagram with server IP and  
port=x; send datagram via  
**clientSocket**

Read datagram from  
**clientSocket**

Close  
**clientSocket**

# Example app: UDP client

## *Python UDPCClient*

```

Include Python's socket library → from socket import *
Create UDP socket for server → clientSocket = socket(AF_INET, SOCK_DGRAM)
Get user keyboard input → message = input('Input lowercase sentence: ')
Attach server name, port to message; send into socket → clientSocket.sendto(message.encode(),
                                                               (serverName, serverPort))
Read reply characters from socket into string → modifiedMessage, serverAddress =
                                                clientSocket.recvfrom(2048)
Print out received string and close socket → print(modifiedMessage.decode())
                                                clientSocket.close()

```

Annotations:

- Include Python's socket library: Points to `from socket import *`
- Create UDP socket for server: Points to `clientSocket = socket(AF_INET, SOCK_DGRAM)`
- Get user keyboard input: Points to `message = input('Input lowercase sentence: ')`
- Attach server name, port to message; send into socket: Points to `clientSocket.sendto(message.encode(), (serverName, serverPort))`
- Read reply characters from socket into string: Points to `modifiedMessage, serverAddress = clientSocket.recvfrom(2048)`
- Print out received string and close socket: Points to `print(modifiedMessage.decode())` and `clientSocket.close()`

Annotations for code elements:

- `serverName`: Points to `serverName = "`
- `serverPort`: Points to `serverPort = 12000`
- `AF_INET`: Points to `clientSocket = socket(AF_INET, SOCK_DGRAM)`
- `SOCK_DGRAM`: Points to `clientSocket = socket(AF_INET, SOCK_DGRAM)`
- `IPv4`: Points to `AF_INET`
- `UDP socket`: Points to `clientSocket = socket(AF_INET, SOCK_DGRAM)`
- `Server IP (Empty=>local system)`: Points to `serverName = "`

# Example app: UDP server

## *Python UDPServer*

```
from socket import *      Listening IP (Empty=>local system)  
serverPort = 12000  
  
Create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)  
  
Bind socket to local port  
number 12000 → serverSocket.bind(("", serverPort))  
print('The server is ready to receive')  
  
Loop forever → while True:  
  
Read from UDP socket into  
message, getting client's  
address (client IP and port) → message, clientAddress = serverSocket.recvfrom(2048)  
modifiedMessage = message.decode().upper()  
  
Send upper case string  
back to this client → serverSocket.sendto(modifiedMessage.encode(),  
                           clientAddress)
```

# Socket programming *with TCP*

---

- Client must contact server
  - Server process must first be running
  - Server must have created socket (door) that welcomes client's contact
  - Client contacts server by:
  - Creating TCP socket, specifying IP address, port number of server process
  - *When client creates socket:* client TCP establishes connection to server TCP
  - When contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
  - Allows server to talk with multiple clients
  - Source port numbers used to distinguish clients.
- Application viewpoint:**
- TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

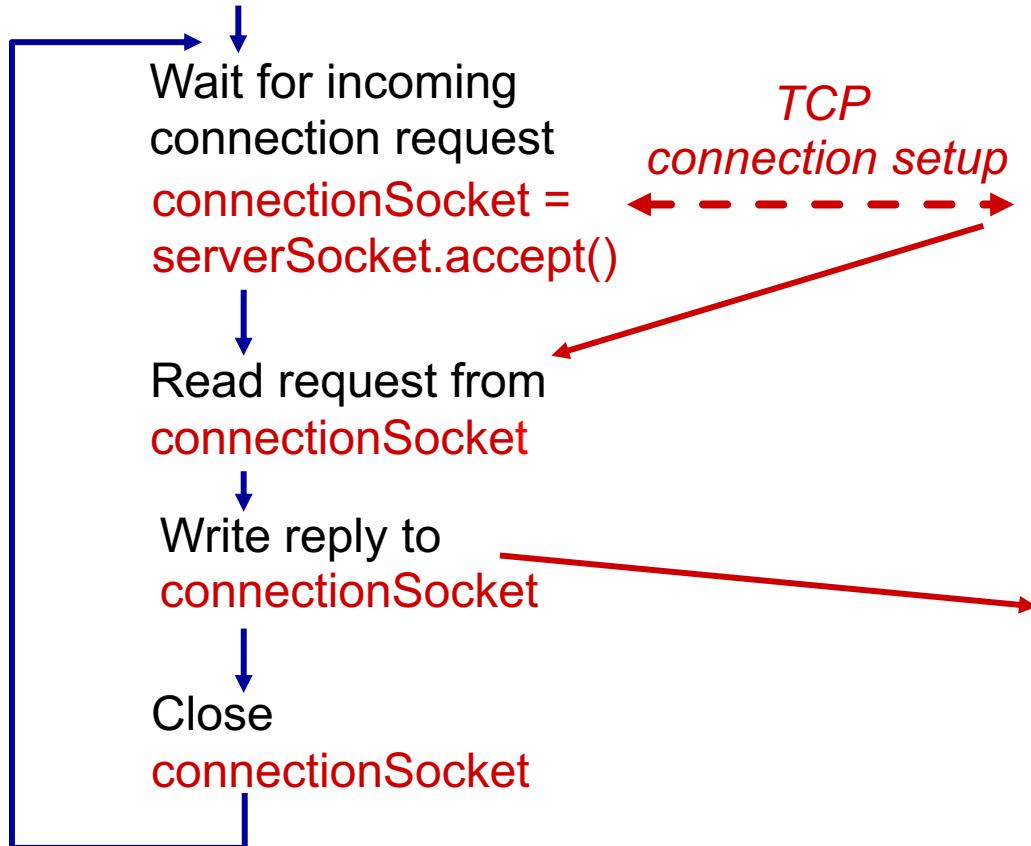
# Client/server socket interaction: TCP

server (running on `hostid`)

Create socket,  
port=**x**, for incoming  
request:

`serverSocket = socket()`

Bind `serverSocket` to an address



Create socket,  
connect to `hostid`, port=**x**  
`clientSocket = socket()`  
`clientSocket.connect()`

Send request using  
`clientSocket`

Read reply from  
`clientSocket`

Close  
`clientSocket`

# Example app: TCP client

We will try to connect to remote port 12000 on the local system

Create a new TCP/IP socket

Attach socket to server name, port

Send data into this socket

Receive data from this socket

Close the socket (ending the connection)

## *Python TCPClient*

```
from socket import *
```

Empty => local machine

IP addressing

Use TCP

```
serverName = ""  
serverPort = 12000
```

```
clientSocket = socket(AF_INET, SOCK_STREAM)
```

```
clientSocket.connect((serverName, serverPort))
```

```
sentence = input('Input lowercase sentence: ')
```

```
clientSocket.send(sentence.encode())
```

```
modifiedSentence = clientSocket.recv(1024)
```

```
print ('From Server: ', modifiedSentence.decode())
```

```
clientSocket.close()
```

# Example app: TCP server

## Python TCP Server

Empty => local machine

Create TCP welcoming socket



Server begins listening for incoming TCP requests



Loop forever



Server waits on accept() for incoming requests, new socket created on return

connectionSocket, addr = serverSocket.accept()

Read bytes from socket (not address as in UDP)

sentence = connectionSocket.recv(1024).decode()

Close connection to this client (but *not* welcoming socket)

capsSentence = sentence.upper()

connectionSocket.send(capsSentence.encode())  
connectionSocket.close()

# Summary

---

- Application architectures
  - Client-server
  - P2P
- Application service requirements
- Internet transport service model
  - Connection-oriented, reliable: TCP
  - Unreliable, datagrams: UDP

Socket programming:  
TCP, UDP sockets

*Important themes:*

- Control vs. messages
  - in-band, out-of-band
- Centralized vs. decentralized
- Stateless vs. stateful
- Reliable vs. unreliable message transfer
- “Complexity at network edge”